

# life\_expectancy

July 3, 2023

```
[313]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn.metrics
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
from statsmodels.tools.tools import add_constant

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.impute import KNNImputer
from sklearn.impute import SimpleImputer
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor
```

This first mini-project is focused on supervised machine learning. For this mini-project I have selected the Life Expectancy dataset that is publicly available on Kaggle.com. From this dataset, I am most interested in understanding the correlation between the factors chosen and longevity, specifically, I am interested in the relative importance of each of these factors in determining longevity. Given my interest is not primarily prediction, but more the predictive power of each factor on longevity, it is important that the resulting model remain interpretable, even at the expense of predictive accuracy.

The life expectancy dataset is a time series over 15 years, with country data from 178 unique countries. contains health and socioeconomic information on over 178 countries. The features of the dataset include: - country (qualitative) - the name of the country the record is relevant for - year (qualitative) - year that the record was taken - status (qualitative) - economic status of the country (Developing/ Developed) - Hepatitis B (quantitative) - percentage of the population that was immunized against Hep B - Measles (quantitative) - cases per 1000 people of Measles - Polio (quantitative) - percentage of the population immunized against Polio - HIV/AIDs (quantitative) - Deaths caused by HIV/AIDs - Infant Deaths (quantitative) - number of infant deaths per 1000 people - under-five deaths (quantitative) - number of deaths of people under 5 years old per 1000 people - total expenditure (quantitative) - The ratio of government medical-health expenses to total government expenses - GDP (quantitative) - Gross Domestic Product - BMI (quantitative) - The average body mass index of the entire population of the country - thinness (quantitative) - Prevalence of thinness among people 19 years old in percentage - Alcohol (quantitative) - Liters of alcohol consumption among people over 15 years old - Schooling (quantitative) - The number of

years that people study - life expectancy (quantitative) - Country life expectancy

There are 2848 total records in the file

```
[314]: df = pd.read_csv('life_expectancy.csv')
df
```

```
[314]:
```

	Country	Year	Status	Population	Hepatitis B	Measles	Polio	
0	Afghanistan	2015	Developing	33736494.0	65.0	1154	6.0	\
1	Afghanistan	2014	Developing	327582.0	62.0	492	58.0	
2	Afghanistan	2013	Developing	31731688.0	64.0	430	62.0	
3	Afghanistan	2012	Developing	3696958.0	67.0	2787	67.0	
4	Afghanistan	2011	Developing	2978599.0	68.0	3013	68.0	
...	...	...	...	...	...	...	...	
2843	Zimbabwe	2004	Developing	12777511.0	68.0	31	67.0	
2844	Zimbabwe	2003	Developing	12633897.0	7.0	998	7.0	
2845	Zimbabwe	2002	Developing	125525.0	73.0	304	73.0	
2846	Zimbabwe	2001	Developing	12366165.0	76.0	529	76.0	
2847	Zimbabwe	2000	Developing	12222251.0	79.0	1483	78.0	
	Diphtheria	HIV/AIDS	infant deaths	under-five deaths				
0	65.0	0.1	62	83	\			
1	62.0	0.1	64	86				
2	64.0	0.1	66	89				
3	67.0	0.1	69	93				
4	68.0	0.1	71	97				
...	...	...	...	...				
2843	65.0	33.6	27	42				
2844	68.0	36.7	26	41				
2845	71.0	39.8	25	40				
2846	75.0	42.1	25	39				
2847	78.0	43.5	24	39				
	Total expenditure	GDP	BMI	thinness	1-19 years	Alcohol		
0	8.16	584.259210	19.1		17.2	0.01	\	
1	8.18	612.696514	18.6		17.5	0.01		
2	8.13	631.744976	18.1		17.7	0.01		
3	8.52	669.959000	17.6		17.9	0.01		
4	7.87	63.537231	17.2		18.2	0.01		
...	...	...	...		...	...		
2843	7.13	454.366654	27.1		9.4	4.36		
2844	6.52	453.351155	26.7		9.8	4.06		
2845	6.53	57.348340	26.3		1.2	4.43		
2846	6.16	548.587312	25.9		1.6	1.72		
2847	7.10	547.358878	25.5		11.0	1.68		
	Schooling	Life expectancy						
0	10.1	65.0						

```

1          10.0          59.9
2           9.9          59.9
3           9.8          59.5
4           9.5          59.2
...
2843        9.2          44.3
2844        9.5          44.5
2845       10.0          44.8
2846        9.8          45.3
2847        9.8          46.0

```

[2848 rows x 18 columns]

```
[315]: df.describe()
```

```

[315]:
      count      Year  Population  Hepatitis B      Measles      Polio  \
mean    2007.500000  1.283457e+07    81.076756    2083.082163    82.682220
std       4.610582  6.196094e+07    25.019068   10249.107207    23.434954
min     2000.000000  3.400000e+01     1.000000     0.000000     3.000000
25%     2003.750000  1.967585e+05    77.000000     0.000000    78.000000
50%     2007.500000  1.391756e+06    92.000000    16.000000    93.000000
75%     2011.250000  7.438947e+06    97.000000   336.750000   97.000000
max     2015.000000  1.293859e+09    99.000000  212183.000000   99.000000

      count  Diphtheria  HIV/AIDS  infant deaths  under-five deaths  \
mean     82.451396    1.756461    28.359902      39.500000
std     23.693936     5.148935   117.188032   159.800866
min       2.000000     0.100000     0.000000     0.000000
25%      78.000000     0.100000     0.000000     0.000000
50%      93.000000     0.100000     3.000000     4.000000
75%      97.000000     0.700000    20.000000    25.000000
max      99.000000    50.600000   1800.000000   2500.000000

      count  Total expenditure      GDP      BMI  thinness  1-19 years  \
mean       5.935577   7664.398813   38.503374      4.847230
std       2.504439  14466.241793   19.955485      4.443695
min       0.370000    1.681350    1.000000     0.100000
25%       4.240000   477.541713   19.500000     1.600000
50%       5.760000  1841.086830   43.900000     3.300000
75%       7.530000  6265.658907   56.200000     7.125000
max      17.600000 119172.741800   77.600000    27.700000

      count  Alcohol  Schooling  Life expectancy
mean    2660.000000  2688.000000    2848.000000

```

mean	4.638932	12.060156	69.347402
std	4.064721	3.320160	9.528332
min	0.010000	0.000000	36.300000
25%	0.930000	10.200000	63.500000
50%	3.785000	12.400000	72.200000
75%	7.810000	14.300000	75.800000
max	17.870000	20.700000	89.000000

## 0.1 Exploratory Data Analysis

In the following, I will go through each feature and perform some initial testing to identify at a high level any relevant insights for each feature.

First, we need to understand how and where there are missing values. Because this is a time series data set, we have the ability to interpolate missing values using an estimator from the existing values. However, before we interpolate, we should consider the distribution of data within each feature, whether it is normalized, and also how many missing values there are and whether those missing values are clustered or normally distributed in the data. There may be some columns from which there is too much data missing for us to realistically impute or interpolate missing values, and for those we will likely populate with the median value from that feature.

Because we are drawing inferences from the data, and our primary goal is to determine which factors contribute to Life Expectancy, with a secondary goal of prediction accuracy, we are going to choose a highly interpretable model. Because the response variable “Life expectancy” is quantitative data, we are going to use linear regression as a model to predict Life expectancy based on the feature set. Linear Regression is highly interpretable, and we will be able to look at the coefficient values in the model summary, as well as the p-values to determine the relative strength and impact of the feature on the model.

## 0.2 First we are going to look at our missing values

```
[316]: #Missing data, first we want to see what data we are missing.
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2848 entries, 0 to 2847
Data columns (total 18 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Country                2848 non-null   object
1   Year                   2848 non-null   int64
2   Status                 2848 non-null   object
3   Population             2204 non-null   float64
4   Hepatitis B            2306 non-null   float64
5   Measles                2848 non-null   int64
6   Polio                  2829 non-null   float64
7   Diphtheria             2829 non-null   float64
8   HIV/AIDS               2848 non-null   float64
9   infant deaths          2848 non-null   int64
```

```

10  under-five deaths      2848 non-null   int64
11  Total expenditure     2627 non-null   float64
12  GDP                   2406 non-null   float64
13  BMI                   2816 non-null   float64
14  thinness 1-19 years   2816 non-null   float64
15  Alcohol               2660 non-null   float64
16  Schooling             2688 non-null   float64
17  Life expectancy       2848 non-null   float64
dtypes: float64(12), int64(4), object(2)
memory usage: 400.6+ KB

```

Above we have all the features and the number of values missing in each feature. It looks like there are many missing values, specifically in Population where we are missing almost 23% of the data. We also know that Population is not normally distributed, and thus likely to be skewed. For this reason, we will likely drop Population data, especially as the impact on Life Expectancy seems questionable.

```

[317]: df_missing = df.set_index(['Country', 'Year'])
df_missing = df_missing[df_missing.isna().any(axis=1)]
print(f'{len(df_missing)/len(df)} % of rows are missing data')

```

```
0.4434691011235955 % of rows are missing data
```


It looks like we're missing quite a lot of data, and that almost half of our records are missing one or more values. Let's take a given feature and see how many countries are missing all of their feature data, vs just some of the feature data.

```

[318]: df_missing_hep = df_missing['Hepatitis B']
df_missing_hep_na = df_missing_hep.isna().groupby(level=0).sum()
df_missing_hep_na[df_missing_hep_na > 10].plot(kind='bar')

```

```
[318]: <Axes: xlabel='Country'>
```



As we can see above, only a handful of countries are missing all of their 'Hepatitis B' data, which means our interpolation strategy of inferring missing feature values based on previous feature values in the time series is promising.

```
[319]: df_pivot = df.pivot(columns='Year',index='Country')
df_pivot.head()
```

```
[319]:
```

	Status						
Year	2000	2001	2002	2003			
Country							
Afghanistan	Developing	Developing	Developing	Developing	\		
Albania	Developing	Developing	Developing	Developing			
Algeria	Developing	Developing	Developing	Developing			
Angola	Developing	Developing	Developing	Developing			
Antigua and Barbuda	Developing	Developing	Developing	Developing			

Year	2004	2005	2006	2007			
Country							
Afghanistan	Developing	Developing	Developing	Developing	\		
Albania	Developing	Developing	Developing	Developing			
Algeria	Developing	Developing	Developing	Developing			
Angola	Developing	Developing	Developing	Developing			
Antigua and Barbuda	Developing	Developing	Developing	Developing			

			... Life expectancy				
Year	2008	2009	...	2006	2007	2008	
Country			...				
Afghanistan	Developing	Developing	...	57.3	57.5	58.1	\
Albania	Developing	Developing	...	74.2	75.9	75.3	
Algeria	Developing	Developing	...	73.4	73.8	74.1	
Angola	Developing	Developing	...	47.7	48.2	48.7	
Antigua and Barbuda	Developing	Developing	...	74.8	75.0	75.2	

Year	2009	2010	2011	2012	2013	2014	2015
Country							
Afghanistan	58.6	58.8	59.2	59.5	59.9	59.9	65.0
Albania	76.1	76.2	76.6	76.9	77.2	77.5	77.8
Algeria	74.4	74.7	74.9	75.1	75.3	75.4	75.6
Angola	49.1	49.6	51.0	56.0	51.1	51.7	52.4
Antigua and Barbuda	75.4	75.6	75.7	75.9	76.1	76.2	76.4

[5 rows x 256 columns]

```
[320]: #Proportion of countries for which there is any missing population data.
print(f' {len(df_pivot.loc[df_pivot.Population.isnull()].any(axis=1))}/len(df.
↪Country.unique())* 100}% of countries are missing population data')
```

23.03370786516854% of countries are missing population data

Below I want to visually represent all of the missing data for Population, just to confirm the

sparseness of the data, before dropping it from the table.

```
[321]: df_pivot = df.pivot(index='Year', columns='Country', values='Population')
std_devs = df_pivot.std().sort_values(ascending=False)

country_order = std_devs.index.tolist()
g = sns.FacetGrid(df, col="Country", col_order=country_order, col_wrap=5,
↳sharex=False)

# Apply a boxplot on each face
g.map(sns.boxplot, "Population")
plt.show()
```

```
/Users/alexcullen/Library/Python/3.10/lib/python/site-
packages/seaborn/axisgrid.py:712: UserWarning: Using the boxplot function
without specifying `order` is likely to produce an incorrect plot.
  warnings.warn(warning)
```





Not only are there many missing values, we can see from the box plots that there are several outliers in the dataset. The data collected for this feature is poor and may bias our dataset.

The only categorical variable in the dataset is economic Status, which is binary and has a value of either “Developed” or “Developing”. We will need to convert this categorical variable into a single dummy variable column with value 1 if the country is “Developed” and 0 otherwise.

Also, looking at the distribution of Status, we can see that every country that was “Developed” at period 0 stayed “Developed” through the last period, and every country that was “Developing” at period 0, remained “Developing” through the last period. This is promising for our interpolation strategy, as the estimator we choose will likely be extremely accurate.


```
[322]: #Economic Status (developing or developed) of each country

df_status = pd.concat([df.drop('Status',axis=1),pd.get_dummies(df.Status)],  
    ↪axis=1)[['Country','Developed']]  
df_status.head()
```

```
[322]:
```


	Country	Developed
0	Afghanistan	False
1	Afghanistan	False
2	Afghanistan	False
3	Afghanistan	False
4	Afghanistan	False

```
[323]: df_status.groupby('Country').sum().hist();
```



Next we are going to take the mean of each feature value by country, and plot those using a box plot to better understand the distribution of data and whether it is normally distributed.

```
[324]: fig, axs = plt.subplots(len(df.drop(['Country', 'Year', 'Status'], axis=1).
    ↪columns), 1, figsize=(10, 5*len(df.drop(['Country', 'Year', 'Status'],
    ↪axis=1).columns)))
for idx, col in enumerate(df.drop(['Country', 'Year', 'Status'], axis=1).
    ↪columns):
    df_pivot = df.pivot(index='Year', columns='Country', values=col)
    means = df_pivot.mean().dropna()
    axs[idx].boxplot(means)
    axs[idx].set_title(col)
    axs[idx].set_ylabel(f'16 year average {col}')
plt.tight_layout()
plt.show()
```




We can see that many of the features are not normally distributed, like GDP, Total Expenditure, HIV/AIDS, and others. These values we will need to log transform so that they do not skew/bias our model's predictions.

Below we look at the correlation between features, and we can see that infant deaths and under-five deaths are almost 1:1 correlated. This will likely be an interaction that we need to account for in our model. Other features are correlated with each other, for example vaccination rates, but not to a material degree that they must be accommodated for.

```
[325]: df_status = pd.concat([df.drop('Status',axis=1),pd.get_dummies(df.Status)],u
    ↪axis=1).drop('Developing',axis=1)
sns.heatmap(df_status.set_index(['Year','Country']).corr())
```

[325]: <Axes: >



In the following section, we are going to perform the following transformations on the data based on what we learned from EDA above. 1. Convert Status as a categorical variable to a numeric

0/1 binary column 2. Log transform positive numeric columns that are not normally distributed 3. Interpolate missing values in features where possible 4. Impute all remaining missing values.

```
[326]: df_melt = df_status.melt(id_vars=['Year', 'Country'],
                               value_vars=df_status.drop(['Year', 'Country'],axis=1).
                               columns,
                               var_name='Feature',
                               value_name='Value')

df_pivot = df_melt.pivot_table(index='Year',
                                columns=['Country', 'Feature'],
                                values='Value')

df_pivot
```

```
[326]: Country Afghanistan
Feature      Alcohol      BMI Developed Diphtheria      GDP HIV/AIDS
Year
2000           0.01    12.2         0.0         24.0      114.56      0.1 \
2001           0.01    12.6         0.0         33.0    117.49698      0.1
2002           0.01    13.0         0.0         36.0    187.84595      0.1
2003           0.01    13.4         0.0         41.0    198.728544      0.1
2004           0.02    13.8         0.0          5.0    219.141353      0.1
2005           0.02    14.2         0.0         58.0     25.29413      0.1
2006           0.03    14.7         0.0         58.0    272.56377      0.1
2007           0.02    15.2         0.0         63.0    369.835796      0.1
2008           0.03    15.7         0.0         64.0    373.361116      0.1
2009           0.01    16.2         0.0         63.0    445.893298      0.1
2010           0.01    16.7         0.0         66.0    553.32894      0.1
2011           0.01    17.2         0.0         68.0    63.537231      0.1
2012           0.01    17.6         0.0         67.0     669.959      0.1
2013           0.01    18.1         0.0         64.0    631.744976      0.1
2014           0.01    18.6         0.0         62.0    612.696514      0.1
2015           0.01    19.1         0.0         65.0    584.25921      0.1

Country      ...      Zimbabwe
Feature Hepatitis B Life expectancy Measles Polio ... Hepatitis B
Year      ...
2000           62.0           54.8    6532.0    24.0 ...      79.0 \
2001           63.0           55.3    8762.0    35.0 ...      76.0
2002           64.0           56.2    2486.0    36.0 ...      73.0
2003           65.0           56.7     798.0    41.0 ...       7.0
2004           67.0           57.0     466.0     5.0 ...      68.0
2005           66.0           57.3    1296.0    58.0 ...      65.0
2006           64.0           57.3    1990.0    58.0 ...      68.0
2007           63.0           57.5    1141.0    63.0 ...      72.0
2008           64.0           58.1    1599.0    64.0 ...      75.0
2009           63.0           58.6    2861.0    63.0 ...      73.0
```

2010	66.0	58.8	1989.0	66.0	...	9.0
2011	68.0	59.2	3013.0	68.0	...	94.0
2012	67.0	59.5	2787.0	67.0	...	97.0
2013	64.0	59.9	430.0	62.0	...	95.0
2014	62.0	59.9	492.0	58.0	...	91.0
2015	65.0	65.0	1154.0	6.0	...	87.0

Country

Feature Life expectancy Measles Polio Population Schooling Total expenditure

Year

2000	46.0	1483.0	78.0	12222251.0	9.8	7.1 \
2001	45.3	529.0	76.0	12366165.0	9.8	6.16
2002	44.8	304.0	73.0	125525.0	10.0	6.53
2003	44.5	998.0	7.0	12633897.0	9.5	6.52
2004	44.3	31.0	67.0	12777511.0	9.2	7.13
2005	44.6	420.0	69.0	129432.0	9.3	6.44
2006	45.4	212.0	71.0	13124267.0	9.5	5.12
2007	46.6	242.0	73.0	1332999.0	9.6	4.47
2008	48.2	0.0	75.0	13558469.0	9.7	4.96
2009	50.0	853.0	69.0	1381599.0	9.9	6.26
2010	52.4	9696.0	89.0	1486317.0	10.0	5.37
2011	54.9	0.0	93.0	14386649.0	10.1	6.31
2012	56.6	0.0	95.0	1471826.0	9.8	6.69
2013	58.0	0.0	95.0	155456.0	10.4	6.88
2014	59.2	0.0	92.0	15411675.0	10.3	6.44
2015	67.0	0.0	88.0	15777451.0	10.3	NaN

Country

Feature infant deaths thinness 1-19 years under-five deaths

Year

2000	24.0	11.0	39.0
2001	25.0	1.6	39.0
2002	25.0	1.2	40.0
2003	26.0	9.8	41.0
2004	27.0	9.4	42.0
2005	28.0	9.0	43.0
2006	28.0	8.6	45.0
2007	29.0	8.2	46.0
2008	30.0	7.8	46.0
2009	30.0	7.5	45.0
2010	29.0	7.1	44.0
2011	28.0	6.8	42.0
2012	26.0	6.5	39.0
2013	25.0	6.2	36.0
2014	23.0	5.9	34.0
2015	22.0	5.6	32.0

[16 rows x 2757 columns]

In this section we are removing all outlier values by country by feature. For example, we look at a given country's rates of vaccination for Hep B, and we remove any extreme (3 std deviations) outliers in Hep B based on the existing data for that country. There are many outlier values in the data set. We will interpolate or impute the missing values that we are removing from the dataset.

```
[327]: threshold = 3

def calculate_z_scores(col):
    if col.std() == 0:
        return [np.nan] * len(col)
    else:
        return (col - col.mean()) / col.std()

z_scores = df_pivot.apply(calculate_z_scores, axis=0)

outliers = np.abs(z_scores) > threshold

df_pivot[outliers] = np.nan
```

```
[328]: df_pre_impute = df_pivot.copy()
```

```
[329]: df_stacked = df_pre_impute.stack(level=0)

df_reset = df_stacked.reset_index()
```

Number of missing values

```
[330]: df_reset.isna().sum().sum()
```

```
[330]: 2612
```

We will use a complex KNN Imputer in order to impute the remaining missing values. This is preferred over a more simple approach because it is likely that there is a high degree of clustering in the dataset. The initial “Status” column is an indication that factors amongst the group of developing or developed nations are dissimilar from each other, but similar within their class.

```
[331]: for col in df_reset.drop('Country',axis =1).columns:
        df_reset[col] = pd.to_numeric(df_reset[col])

numeric_columns = df_reset.select_dtypes(include=[np.number])

imputer = KNNImputer(n_neighbors=5)

df_imputed_numeric = imputer.fit_transform(numeric_columns)

df_imputed_numeric = pd.DataFrame(df_imputed_numeric, columns=numeric_columns.
    ↪columns)
```



```
df_imputed = pd.concat([df_reset.drop(numeric_columns.columns, axis=1),
↳ df_imputed_numeric], axis=1)
```

Here we are log transforming the non-normally distributed columns

```
[332]: df_transformed = df_imputed.copy()

log_transform_cols = ['GDP', 'Total expenditure', 'under-five deaths', 'infant_
↳ deaths', 'Measles', 'HIV/AIDS']

for col in log_transform_cols:
    df_transformed[col] = df_transformed[col].apply(lambda x: np.log(x+ 1))
```

```
[333]: df_transformed
```

```
[333]: Feature          Country  Year  Alcohol  BMI  Developed \
0          Afghanistan  2000.0    0.010  12.2         0.0 \
1              Albania  2000.0    3.660  45.0         0.0
2              Algeria  2000.0    0.250  44.4         0.0
3              Angola  2000.0    1.850  15.4         0.0
4  Antigua and Barbuda  2000.0    7.270  38.2         0.0
...
2843  Venezuela (Bolivarian Republic of)  2015.0    5.784  62.1         0.0
2844              Viet Nam  2015.0    1.248  17.5         0.0
2845              Yemen  2015.0    2.444  41.3         0.0
2846              Zambia  2015.0    3.048  23.4         0.0
2847              Zimbabwe  2015.0    1.582  31.8         0.0
```

```
Feature  Diphtheria      GDP  HIV/AIDS  Hepatitis B  Life expectancy \
0          24.0  4.749790  0.095310          62.0          54.8 \
1          97.0  7.070545  0.095310          96.0          72.6
2          86.0  7.472033  0.095310          86.0          71.3
3          28.0  6.321302  1.098612          69.0          45.3
4          95.0  9.197879  0.095310          96.4          73.6
...
2843          87.0  7.947697  0.095310          87.0          74.1
2844          97.0  6.280968  0.095310          97.0          76.0
2845          69.0  8.116740  0.095310          69.0          65.7
2846           9.0  7.181508  1.629241           9.0          61.8
2847          87.0  4.784937  1.974081          87.0          67.0
```

```
Feature  Measles  Polio    Population  Schooling  Total expenditure \
0      8.784622   24.0  2.937560e+05         5.5      2.219203 \
1      2.104134   97.0  3.892700e+04        10.7      1.982380
2      0.000000   86.0  3.118366e+06        10.7      1.501853
3      7.705262    3.0  1.644924e+06         4.6      1.332366
4      0.000000   96.0  1.058848e+07         0.0      1.635106
```

...	...	...	...	...	...
2843	0.000000	87.0	2.258149e+07	14.3	1.926873
2844	5.549076	97.0	2.354233e+06	12.6	1.746762
2845	6.150603	63.0	1.281210e+07	9.0	1.781036
2846	2.302585	9.0	1.615870e+05	12.5	1.796083
2847	0.000000	88.0	1.577745e+07	10.3	1.907466

Feature	infant deaths	thinness	1-19 years	under-five deaths
0	4.488636		2.3	4.812184
1	0.693147		2.1	0.693147
2	3.091042		6.5	3.258097
3	4.584967		1.9	5.081404
4	0.000000		3.7	0.000000

...	...	...	...
2843	2.302585	1.6	2.397895
2844	3.367296	14.2	3.583519
2845	3.637586	13.6	3.871201
2846	3.332205	6.3	3.713572
2847	3.135494	5.6	3.496508

[2848 rows x 18 columns]

Now that all of the data is cleaning and transformed, we will run a simple OLS linear regrssion using Statsmodel as our baseline to see how a normal linear regression model performs.

```
[334]: df_final = df_transformed.set_index(['Year', 'Country']).
        drop('Population',axis=1)
X = df_final.drop('Life expectancy', axis=1)
y = df_final['Life expectancy']

X = sm.add_constant(X)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
        random_state=42)

model = sm.OLS(y_train, X_train)
results = model.fit()

print(results.summary())
```

#### OLS Regression Results

Dep. Variable:	Life expectancy	R-squared:	0.847
Model:	OLS	Adj. R-squared:	0.846
Method:	Least Squares	F-statistic:	897.8
Date:	Mon, 03 Jul 2023	Prob (F-statistic):	0.00
Time:	21:22:55	Log-Likelihood:	-6218.5
No. Observations:	2278	AIC:	1.247e+04

```

Df Residuals:          2263    BIC:          1.255e+04
Df Model:              14
Covariance Type:      nonrobust
=====
=====
              coef      std err          t      P>|t|      [0.025
0.975]
-----
-----
const          54.3579      0.773      70.338      0.000      52.842
55.873
Alcohol        -0.0366      0.028      -1.324      0.186      -0.091
0.018
BMI             0.0207      0.006       3.716      0.000       0.010
0.032
Developed       2.2710      0.277       8.184      0.000       1.727
2.815
Diphtheria      0.0340      0.006       6.151      0.000       0.023
0.045
GDP             0.4779      0.057       8.445      0.000       0.367
0.589
HIV/AIDS        -5.5321      0.130     -42.499      0.000     -5.787
-5.277
Hepatitis B     -0.0155      0.004      -3.543      0.000     -0.024
-0.007
Measles         0.0257      0.033       0.767      0.443     -0.040
0.091
Polio           0.0207      0.005       3.888      0.000       0.010
0.031
Schooling       0.8868      0.040     22.346      0.000       0.809
0.965
Total expenditure 0.4077      0.216       1.887      0.059     -0.016
0.831
infant deaths    3.1913      0.579       5.509      0.000       2.055
4.327
thinness 1-19 years -0.0517      0.023      -2.203      0.028     -0.098
-0.006
under-five deaths -3.6897      0.555     -6.651      0.000     -4.778
-2.602
=====
Omnibus:          74.006    Durbin-Watson:          2.005
Prob(Omnibus):    0.000    Jarque-Bera (JB):      159.576
Skew:            -0.185    Prob(JB):              2.23e-35
Kurtosis:         4.243    Cond. No.              1.71e+03
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly

specified.

[2] The condition number is large,  $1.71e+03$ . This might indicate that there are strong multicollinearity or other numerical problems.

We can see from the summary table that Simple Linear Regression is actually quite effective at modeling the data. The adj R-squared for our model is 84.6% which means that our model explains 84.6% of the variance in the errors.

Looking at each parameter and the assessment, we can see based on the coefficient and p-value of each the relevance and impact of each parameter on the model. Immediately we can see that Total Expenditure on healthcare is above our p-value threshold of .05, as well as # of cases of measles. We can see that infant deaths and under-five deaths are extremely correlated with each other, and so we will need to include that interaction in our model.

From first review, it seems that HIV/AIDS is the strongest predictor of life expectancy, followed by whether or not the country is “Developed”.

```
[335]: y_pred = results.predict(X_test)

mean_squared_error(y_test, y_pred) #baseline score
```

```
[335]: 13.256831860025624
```

To better analyze multicollinearity, we will use VIF below. We note that infant deaths and under-five deaths are highly correlated.

```
[336]: vif_data = pd.DataFrame()
vif_data["feature"] = X.columns

vif_data["VIF"] = [variance_inflation_factor(X.values, i) for i in range(len(X.
↪columns))]

print(vif_data)
```

	feature	VIF
0	const	99.328038
1	Alcohol	1.985335
2	BMI	1.919887
3	Developed	1.822120
4	Diphtheria	2.472871
5	GDP	1.701669
6	HIV/AIDS	1.585255
7	Hepatitis B	1.629814
8	Measles	1.824198
9	Polio	2.145212
10	Schooling	2.752249
11	Total expenditure	1.141331
12	infant deaths	145.658011
13	thinness 1-19 years	1.775081
14	under-five deaths	153.457055

Here we are creating the interaction between infant deaths and under-five deaths

```
[337]: df_forward_select = df_final.copy()

df_forward_select['infant death * under-five deaths'] =
↳df_forward_select['infant deaths'] * df_forward_select['under-five deaths']
```

We split the data to validate the accuracy of our predictions on unseen data.

```
[338]: X = df_forward_select.drop('Life expectancy', axis=1)
y = df_forward_select['Life expectancy']

X = sm.add_constant(X)

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

Below we are going to implement the hybrid approach of forward selecting and backwards selecting features in each iteration of the loop until our model settles and we are no longer adding features. Each forward step we are adding the feature which improves the mean\_squared\_error the most, each backward step we are removing any features with p-values > our threshold.

```
[339]: def get_score(X_train, X_test, y_train, y_test):
    model = sm.OLS(y_train, add_constant(X_train))
    results = model.fit()
    y_pred = results.predict(add_constant(X_test))
    return mean_squared_error(y_test, y_pred)

remaining = set(X_train.columns)
selected = []
n = len(X_train)
best_new_score, current_score = float('inf'), float('inf')
while remaining:
    changed = False

    # Forward step
    scores_with_candidates = []
    for candidate in remaining:
        temp_selected = selected + [candidate]
        score = get_score(X_train[temp_selected], X_test[temp_selected],
↳y_train, y_test)
        scores_with_candidates.append((score, candidate))

    scores_with_candidates.sort()
    best_new_score, best_candidate = scores_with_candidates[0]

    if best_new_score < current_score:
        remaining.remove(best_candidate)
```

```

        selected.append(best_candidate)
        current_score = best_new_score
        changed = True

    # Backward step
    model = sm.OLS(y_train, add_constant(X_train[selected]))
    pvalues = model.fit().pvalues[1:] # exclude intercept
    worst_pval = pvalues.max()
    if worst_pval > 0.05:
        worst_feature = pvalues.idxmax()
        selected.remove(worst_feature)
        remaining.add(worst_feature)
        changed = True

    if not changed:
        break

best_model = sm.OLS(y_train, add_constant(X_train[selected]))
results = best_model.fit()
print(results.summary())

```

#### OLS Regression Results

```

=====
Dep. Variable:          Life expectancy    R-squared:                0.846
Model:                  OLS               Adj. R-squared:           0.845
Method:                 Least Squares     F-statistic:             1246.
Date:                   Mon, 03 Jul 2023   Prob (F-statistic):       0.00
Time:                   21:22:56          Log-Likelihood:          -6228.6
No. Observations:      2278              AIC:                    1.248e+04
Df Residuals:           2267              BIC:                    1.254e+04
Df Model:               10
Covariance Type:        nonrobust
=====

```

```

=====

```

	coef	std err	t	P> t	[0.025
0.975]					
-----					
const	55.5427	0.683	81.315	0.000	54.203
56.882					
HIV/AIDS	-5.6245	0.126	-44.570	0.000	-5.872
-5.377					
Schooling	0.9070	0.038	24.127	0.000	0.833
0.981					
GDP	0.4986	0.056	8.942	0.000	0.389
0.608					
under-five deaths	-3.8362	0.555	-6.917	0.000	-4.924
-2.749					

Polio	0.0207	0.005	3.884	0.000	0.010
0.031					
Developed	2.1944	0.251	8.755	0.000	1.703
2.686					
infant deaths	3.3615	0.580	5.795	0.000	2.224
4.499					
Diphtheria	0.0347	0.006	6.261	0.000	0.024
0.046					
Hepatitis B	-0.0154	0.004	-3.523	0.000	-0.024
-0.007					
thinness 1-19 years	-0.0767	0.022	-3.550	0.000	-0.119
-0.034					

Omnibus:	70.443	Durbin-Watson:	1.998
Prob(Omnibus):	0.000	Jarque-Bera (JB):	142.079
Skew:	-0.198	Prob(JB):	1.41e-31
Kurtosis:	4.157	Cond. No.	1.58e+03

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.58e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[340]: print(f'Final score = {mean_squared_error(y_test, results.
        ↪predict(add_constant(X_test[selected]))})')
print(f'Selected attributes ranked in order of adj. r-squared ')
[col for col in selected]
```

Final score = 13.154045372207346

Selected attributes ranked in order of adj. r-squared

```
[340]: ['HIV/AIDS',
        'Schooling',
        'GDP',
        'under-five deaths',
        'Polio',
        'Developed',
        'infant deaths',
        'Diphtheria',
        'Hepatitis B',
        'thinness 1-19 years']
```

## 1 Final Analysis:

Based on our modeling, it is clear that the most immediate predictor of a country's Life Expectancy is it's incidence of HIV/AIDS, which is a deadly and serious disease. The correlation is strongly

negative which means that the higher the prevalence the lower the Life Expectancy of that country.

After that, we see that metrics related to wealth like Schooling, GDP, and Status are positively correlated with Life Expectancy, which is not unusual. However, what is unusual is that Total expenditure on health care does not at all have predictive power on life expectancy.

Linear Regression was a very effective model for the analysis that we were performing today. It is surprising to see how accurate our model was, with our best performing model able to predict a given country's life expectancy with a root mean squared error of 4 years, which is a highly accurate measurement given the range of Life expectancies seen in the data.

Also, with Linear Regression we did not compromise on interpretability, we were readily able to infer from the data which factors were most important in predicting Life Expectancy by country

```
[354]: #As a final step, we will look at the relative performance of KNN which is more
      ↪difficult to interpret, but might be more accurate at predicting life
      ↪expectancy.
```

```
[341]: model = KNeighborsRegressor(n_neighbors=3)
      model.fit(X_train, y_train)

      y_pred = model.predict(X_test)

      mean_squared_error(y_test, y_pred)
```

```
[341]: 15.671372553606236
```

```
[356]: X_train.loc[2000.0]
```

```
[356]:
```

	const	Alcohol	BMI	Developed	Diphtheria	GDP	HIV/AIDS
Country							
Timor-Leste	1.0	0.50	11.9	0.0	93.6	6.048049	0.095310 \
Mauritius	1.0	4.60	25.3	0.0	88.0	8.259024	0.095310
Gambia	1.0	2.18	18.0	0.0	85.6	5.088505	1.098612
Hungary	1.0	12.22	56.1	1.0	99.0	8.439116	0.095310
Israel	1.0	2.53	58.3	0.0	93.0	7.674684	0.095310
...	...	...	...	...	...	...	
Guinea	1.0	0.17	16.6	0.0	46.0	3.570009	1.386294
Tunisia	1.0	1.21	48.1	0.0	97.0	7.702969	0.095310
Comoros	1.0	0.09	17.3	0.0	7.0	5.931855	0.095310
Botswana	1.0	5.37	29.9	0.0	97.0	8.116921	3.683867
Saint Lucia	1.0	11.69	36.8	0.0	7.0	7.553143	0.336472

	Hepatitis B	Measles	Polio	Schooling	Total expenditure
Country					
Timor-Leste	94.6	0.000000	90.8	7.68	1.449269 \
Mauritius	88.0	0.000000	88.0	12.10	1.564441
Gambia	91.0	5.519860	84.0	6.50	1.528228
Hungary	83.2	0.693147	99.0	13.90	2.151762



Israel	98.0	3.610918	93.0	15.20	2.095561
...	...	...	...	...	...
Guinea	36.2	9.067809	47.0	4.80	1.495149
Tunisia	94.0	3.871201	97.0	12.80	1.856298
Comoros	65.8	0.000000	7.0	8.20	1.517323
Botswana	86.0	0.336472	97.0	11.70	1.729884
Saint Lucia	42.6	0.000000	7.0	12.80	1.876407

	infant deaths	thinness	1-19 years	under-five deaths
Country				
Timor-Leste	1.386294		12.2	1.609438 \
Mauritius	0.000000		8.1	0.000000
Gambia	1.386294		1.2	1.945910
Hungary	0.693147		2.3	0.693147
Israel	0.693147		1.1	0.693147
...	...		...	...
Guinea	3.637586		1.3	4.110874
Tunisia	1.609438		6.6	1.791759
Comoros	0.693147		7.9	1.098612
Botswana	1.098612		12.3	1.609438
Saint Lucia	0.000000		4.5	0.000000

	infant death * under-five deaths
Country	
Timor-Leste	2.231155
Mauritius	0.000000
Gambia	2.697604
Hungary	0.480453
Israel	0.480453
...	...
Guinea	14.953658
Tunisia	2.883726
Comoros	0.761500
Botswana	1.768148
Saint Lucia	0.000000

[146 rows x 16 columns]

```
[379]: X_test.head()
```

```
[379]:
```

		const	Alcohol	BMI	Developed	Diphtheria	GDP
Year	Country						
2008.0	Switzerland	1.0	10.29	54.6	1.0	95.0	11.186095 \
2001.0	Spain	1.0	9.86	58.2	1.0	96.0	9.637215
2011.0	Ukraine	1.0	8.48	59.0	0.0	5.0	8.180533
2014.0	Brazil	1.0	7.32	55.3	0.0	93.0	7.112830
2012.0	Cyprus	1.0	10.55	58.7	1.0	99.0	10.273400

Year	Country	HIV/AIDS	Hepatitis B	Measles	Polio	Schooling
2008.0	Switzerland	0.095310	98.0	7.612337	96.0	15.3 \
2001.0	Spain	0.095310	83.0	0.000000	95.0	15.7
2011.0	Ukraine	0.182322	21.0	7.195937	54.0	14.9
2014.0	Brazil	0.095310	96.0	5.564520	96.0	15.2
2012.0	Cyprus	0.095310	96.0	0.693147	99.0	13.8

Year	Country	Total expenditure	infant deaths	thinness	1-19 years
2008.0	Switzerland	0.828552	0.000000		0.5 \
2001.0	Spain	2.109000	1.098612		0.6
2011.0	Ukraine	2.076938	1.791759		2.4
2014.0	Brazil	2.232163	3.806662		2.7
2012.0	Cyprus	2.132982	0.000000		0.9

Year	Country	under-five deaths	infant death * under-five deaths
2008.0	Switzerland	0.000000	0.000000
2001.0	Spain	1.098612	1.206949
2011.0	Ukraine	1.791759	3.210402
2014.0	Brazil	3.912023	14.891751
2012.0	Cyprus	0.000000	0.000000

```
[383]: distances, indices = model.kneighbors(X=pd.DataFrame(X_test.loc[2008.0].
↳loc['Switzerland']).T)
```

```
[385]: print(distances)
X_test.iloc[indices[0]]
```

```
[[1.86587358 2.40112778 3.79775239]]
```

```
[385]:
```

Year	Country	const	Alcohol	BMI	Developed	Diphtheria	GDP
2005.0	Ireland	1.0	13.31	55.1	1.0	96.0	8.680642 \
2003.0	Israel	1.0	2.32	59.6	0.0	93.0	9.849454
2004.0	Senegal	1.0	0.35	19.1	0.0	87.0	6.599301

Year	Country	HIV/AIDS	Hepatitis B	Measles	Polio	Schooling
2005.0	Ireland	0.095310	98.4	4.564348	96.8	17.5 \
2003.0	Israel	0.095310	98.0	4.828314	93.0	16.0
2004.0	Senegal	0.530628	76.0	3.465736	87.0	6.2

Year	Country	Total expenditure	infant deaths	thinness	1-19 years
2005.0	Ireland	2.112635	0.000000		0.3 \

2003.0	Israel	2.132982	0.693147	1.1
2004.0	Senegal	1.899118	3.218876	11.6

Year	Country	under-five deaths	infant death * under-five deaths
2005.0	Ireland	0.000000	0.000000
2003.0	Israel	0.693147	0.480453
2004.0	Senegal	3.761200	12.106836

```
[382]: y_test.iloc[indices[0]]
```

```
[382]: Year    Country
2005.0  Ireland    78.7
2003.0   Israel    79.7
2004.0  Senegal    59.7
Name: Life expectancy, dtype: float64
```

Above we performed KNN on the same transformed and cleaned dataset, and we achieved a lower accuracy score than even the baseline Linear Regression model. We can also see that when use KNN to return the nearest neighbors, we get some non-intuitive results upon further analysis. For example, 2008 Switzerland's nearest neighbors are Ireland and Israel which make sense, but then Senegal is the third closest, however, as we can see from the data, the Life Expectancy of the Senegalese is far below Switzerland's own.

```
[ ]:
```